

DS HW1

Ruihan Ding

2026-02-25

```
library(ISLR)
library(glmnet)
library(caret)
library(tidymodels)
library(corrplot)
library(ggplot2)
library(plotmo)
library(ggrepel)
library(tidyverse)
library(pls)
```

Import the data.

```
housing_train =
  read_csv("housing_training.csv") |>
  janitor::clean_names()

## Rows: 1440 Columns: 26
## -- Column specification -----
## Delimiter: ","
## chr (4): Overall_Qual, Kitchen_Qual, Fireplace_Qu, Exter_Qual
## dbl (22): Gr_Liv_Area, First_Flr_SF, Second_Flr_SF, Total_Bsmt_SF, Low_Qual...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

housing_test =
  read_csv("housing_test.csv") |>
  janitor::clean_names()

## Rows: 959 Columns: 26
## -- Column specification -----
## Delimiter: ","
## chr (4): Overall_Qual, Kitchen_Qual, Fireplace_Qu, Exter_Qual
## dbl (22): Gr_Liv_Area, First_Flr_SF, Second_Flr_SF, Total_Bsmt_SF, Low_Qual...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

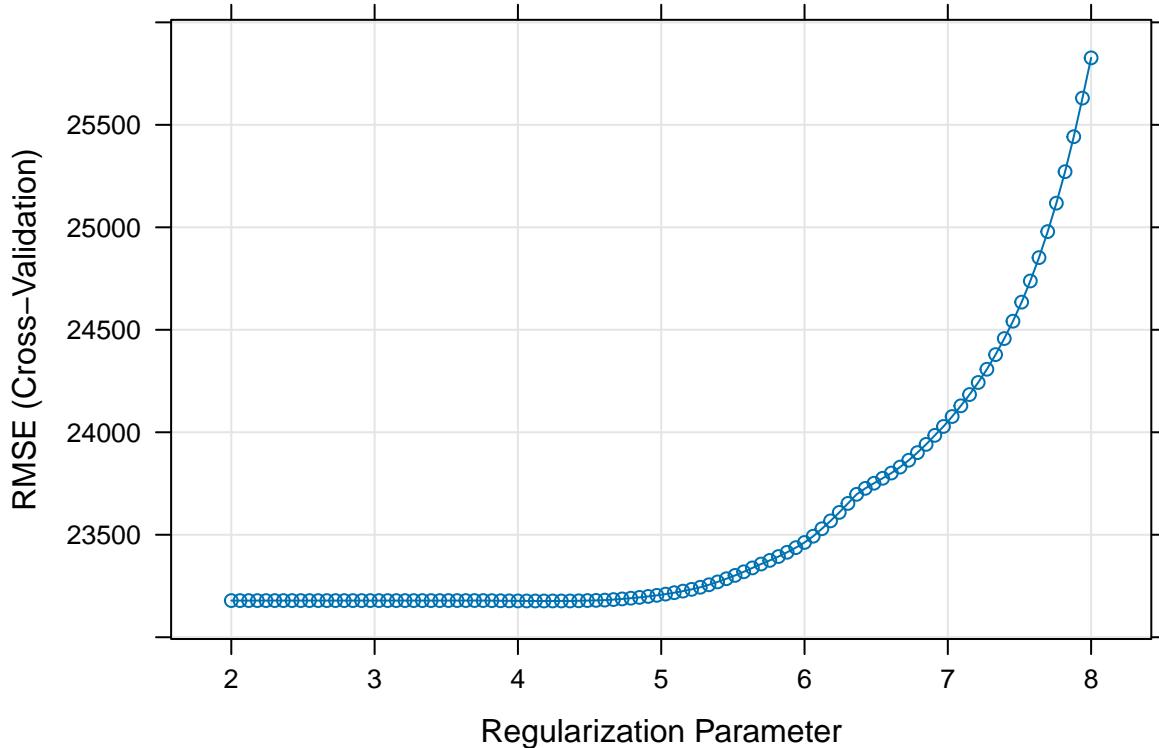
a

- Fit a lasso model on the training data.

```
ctrl1 = trainControl(method = "cv",
                     number = 10,
                     selectionFunction = "best")

set.seed(2026)
lasso.fit = train(sale_price ~ .,
                  data = housing_train,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = 1,
                                         lambda = exp(seq(8, 2, length = 100))),
                  trControl = ctrl1)

plot(lasso.fit, xTrans = log)
```



- Report the selected tuning parameter and the test error.

```
lasso_lambda_min = lasso.fit$bestTune$lambda

lasso.pred = predict(lasso.fit, newdata = housing_test)
lasso_test_mse = mean((lasso.pred - housing_test[["sale_price"]])^2)
```

The selected lambda is 65.4848083, and the test error is 4.3996792×10^8 .

- When the 1SE rule is applied, how many predictors are included in the model?

```

ctrl_1se = trainControl(method = "cv",
                        number = 10,
                        selectionFunction = "oneSE")

set.seed(2026)
lasso_1se.fit = train(sale_price ~ .,
                      data = housing_train,
                      method = "glmnet",
                      tuneGrid = expand.grid(alpha = 1,
                                             lambda = exp(seq(8, 2, length = 100))),
                      trControl = ctrl_1se)

lasso_1se = lasso_1se.fit$bestTune$lambda
lasso_1se_coef = coef(lasso_1se.fit$finalModel, lasso_1se)
lasso_1se_coef

```

```

## 40 x 1 sparse Matrix of class "dgCMatrix"
##                                     s=942.4519
## (Intercept)                 -1.708349e+06
## gr_liv_area                  5.568317e+01
## first_flr_sf                  1.178647e+00
## second_flr_sf                   .
## total_bsmt_sf                  3.675124e+01
## low_qual_fin_sf                -2.390210e+01
## wood_deck_sf                  8.085711e+00
## open_porch_sf                  7.140424e+00
## bsmt_unf_sf                   -1.899986e+01
## mas_vnr_area                  1.436778e+01
## garage_cars                   3.087641e+03
## garage_area                    1.159899e+01
## year_built                     3.169102e+02
## tot_rms_abv_grd                -8.706872e+02
## full_bath                      .
## overall_qualAverage            -2.900389e+03
## overall_qualBelow_Average     -8.672363e+03
## overall_qualExcellent          8.933970e+04
## overall_qualFair               -5.488845e+03
## overall_qualGood                9.437591e+03
## overall_qualVery_Excellent    1.587474e+05
## overall_qualVery_Good          3.558537e+04
## kitchen_qualFair                -4.549109e+03
## kitchen_qualGood                  .
## kitchen_qualTypical            -9.766031e+03
## fireplaces                      6.635631e+03
## fireplace_quFair                  .
## fireplace_quGood                 4.494381e+03
## fireplace_quNo_Fireplace        .
## fireplace_quPoor                  .
## fireplace_quTypical                 .
## exter_qualFair                  -1.387356e+04
## exter_qualGood                  .
## exter_qualTypical                -5.269069e+03
## lot_frontage                     6.524291e+01

```

```

## lot_area           5.461956e-01
## longitude        -6.992835e+03
## latitude          1.127163e+04
## misc_val           .
## year_sold          .

```

Under the 1SE rule, the model includes 29 predictors.

b

- Fit an elastic net model on the training data.

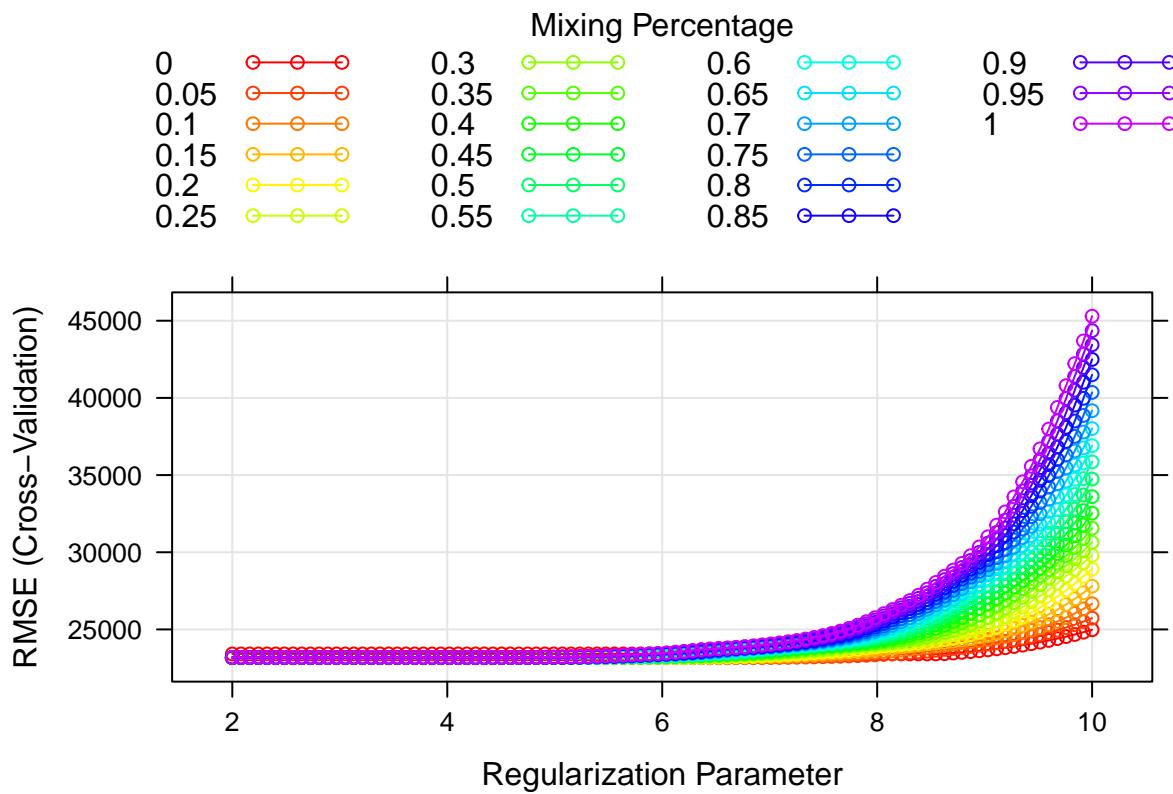
```

set.seed(2026)
enet.fit = train(sale_price ~ .,
                  data = housing_train,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                         lambda = exp(seq(10, 2, length = 100))),
                  trControl = ctrl1)

myCol = rainbow(25)
myPar = list(superpose.symbol = list(col = myCol),
             superpose.line = list(col = myCol))

plot(enet.fit, par.settings = myPar, xTrans = log)

```



- Report the selected tuning parameters and the test error.

```

enet.fit$bestTune

##      alpha lambda
## 156  0.05 629.197

enet_alpha = enet.fit$bestTune$alpha
enet_lambda_min = enet.fit$bestTune$lambda

enet.pred = predict(enet.fit, newdata = housing_test)
enet_test_mse = mean((enet.pred - housing_test[["sale_price"]])^2)

```

The selected alpha is 0.05, lambda is 629.1970259, and the test error is 4.381072×10^8 .

- Is it possible to apply the 1SE rule to select the tuning parameters for elastic net? If the 1SE rule is applicable, implement it to select the tuning parameters. If not, explain why.

```

set.seed(2026)
enet_1se.fit = train(sale_price ~ .,
                      data = housing_train,
                      method = "glmnet",
                      tuneGrid = expand.grid(alpha = enet_alpha,
                                             lambda = exp(seq(10, 2, length = 100))),
                      trControl = ctrl_1se)

enet_1se = enet_1se.fit$bestTune$lambda

```

Based on the previously selected optimal alpha = 0.05, I applied the 1SE rule to determine lambda, which is 7105.9438067.

C

- Fit a partial least squares model on the training data and report the test error.

```

x = model.matrix(sale_price ~ ., housing_train)[, -1]
y = housing_train[["sale_price"]]

x2 = model.matrix(sale_price ~ ., housing_test)[, -1]
y2 = housing_test[["sale_price"]]

set.seed(2026)
pls_fit = train(x, y,
                 method = "pls",
                 tuneGrid = data.frame(ncomp = 1:19),
                 trControl = ctrl1,
                 preProcess = c("center", "scale"))

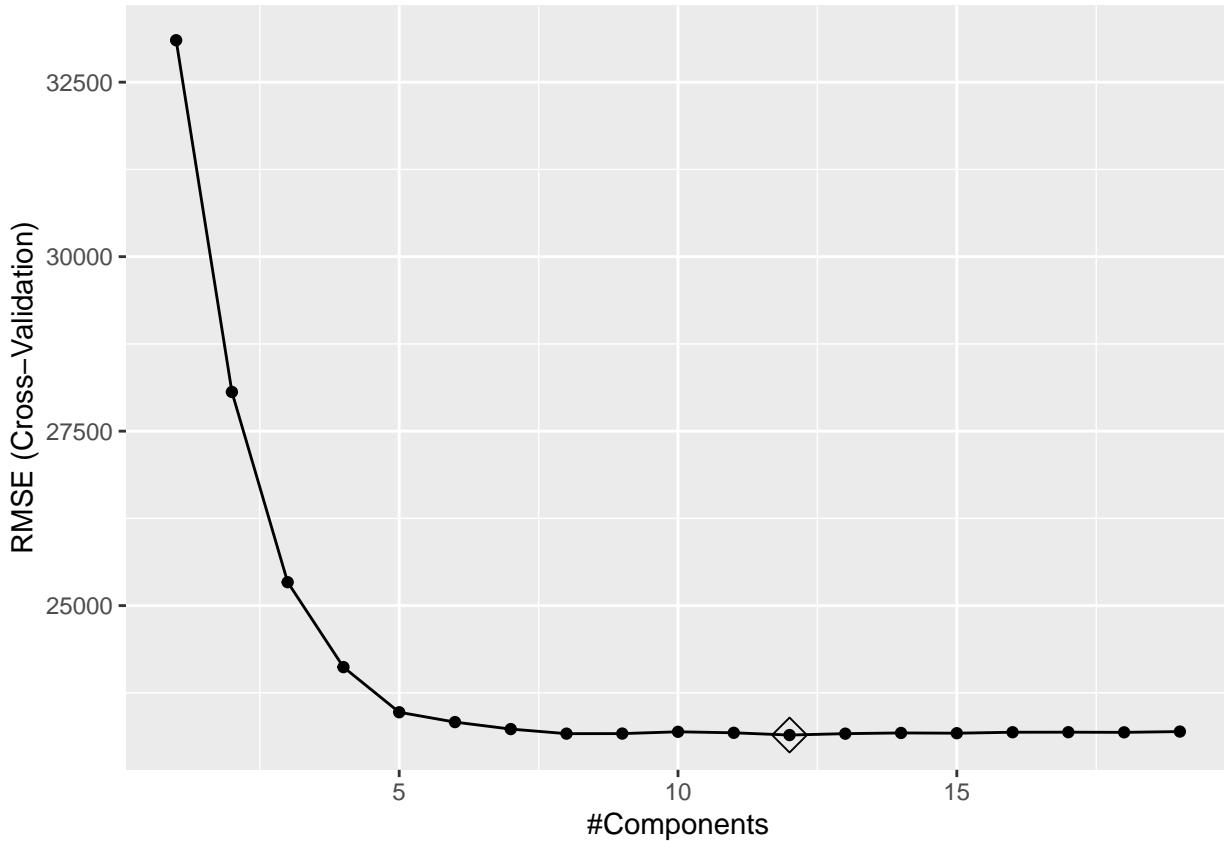
predy_pls = predict(pls_fit, newdata = x2)
pls_test_mse = mean((y2 - predy_pls)^2)

```

The test error is 4.4962272×10^8 .

- How many components are included in your model?

```
ggplot(pls_fit, highlight = TRUE)
```



There are 12 components in my model.

d

- Choose the best model for predicting the response and explain your choice.

```
resamp = resamples(list(lasso = lasso.fit, lasso_1se = lasso_1se.fit, enet = enet.fit, enet_1se = enet_1se.fit))

summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lasso, lasso_1se, enet, enet_1se, pls
## Number of resamples: 10
##
## MAE
##      Min. 1st Qu. Median     Mean 3rd Qu.    Max. NA's
## lasso   14767.34 16239.16 16549.68 16668.63 17386.11 18238.66  0
```

```

## lasso_1se 15033.07 16138.48 16639.94 16885.36 17969.81 18574.42      0
## enet       14738.00 16197.92 16513.52 16639.35 17367.60 18182.56      0
## enet_1se   14874.30 16280.09 16636.11 16818.23 17745.68 18166.41      0
## pls        14893.35 16234.26 16664.76 16712.82 17444.65 18244.49      0
##
## RMSE
##           Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## lasso      20061.98 21746.27 22975.06 23176.81 24580.80 28303.85      0
## lasso_1se  20664.02 21951.74 23537.46 23940.75 25205.61 30682.77      0
## enet       20039.62 21714.28 23015.51 23166.35 24604.52 28258.98      0
## enet_1se   20815.36 22069.64 23505.00 23900.49 24661.24 30952.67      0
## pls        20051.99 21777.34 23017.97 23145.53 24573.40 27804.33      0
##
## Rsquared
##           Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## lasso      0.8705812 0.8851274 0.9075688 0.9016793 0.9169088 0.9263271      0
## lasso_1se  0.8667812 0.8758247 0.9030785 0.8954890 0.9098742 0.9258259      0
## enet       0.8707239 0.8849025 0.9075583 0.9018122 0.9167978 0.9261045      0
## enet_1se   0.8691088 0.8777841 0.9015654 0.8966101 0.9108552 0.9239632      0
## pls        0.8693941 0.8855808 0.9074139 0.9018073 0.9168251 0.9257025      0

```

Partial least squares model is the best choice for predicting the response, because it achieves the lowest cross-validated RMSE.

e

- If R package “caret” was used for the lasso in (a), retrain this model using R package “glmnet”, and vice versa.

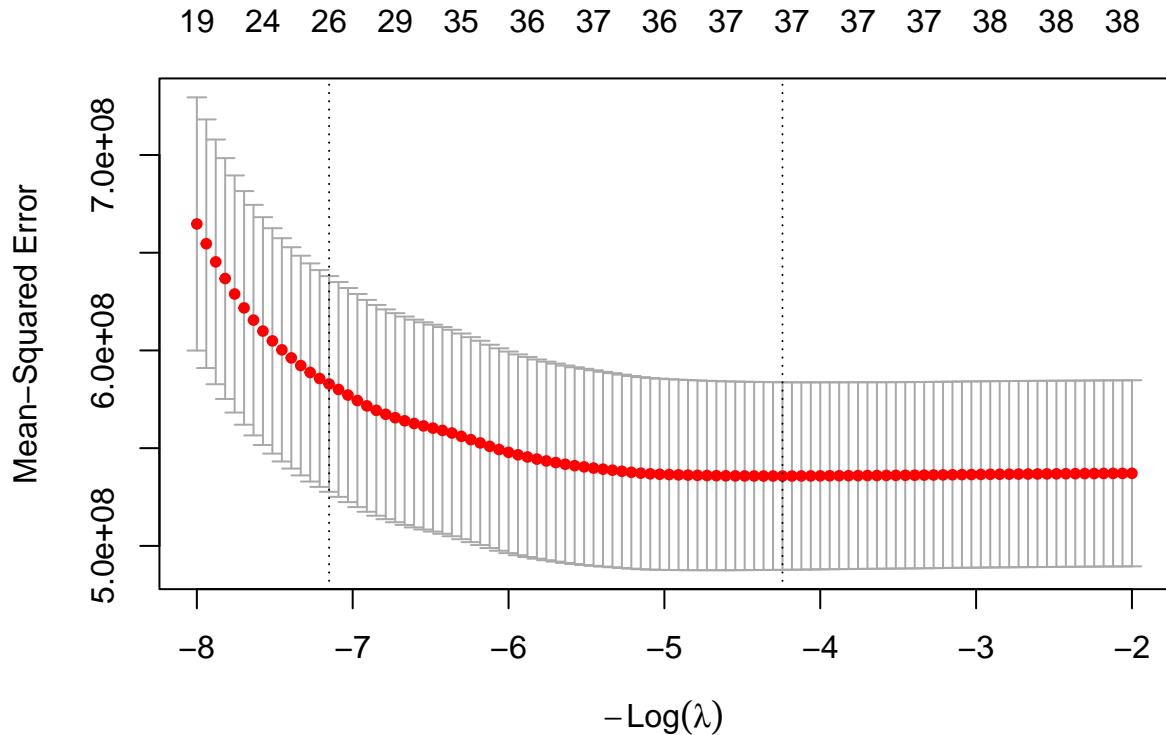
```

set.seed(2026)
cv.lasso = cv.glmnet(x, y,
                      alpha = 1,
                      lambda = exp(seq(8, 2, length = 100)))

```

- Compare the selected tuning parameters between the two software approaches. Should there be discrepancies in the chosen parameters, discuss potential reasons for these differences.

```
plot(cv.lasso)
```



```
cv.lasso$lambda.min
```

```
## [1] 69.57632
```

```
cv.lasso$lambda.1se
```

```
## [1] 1276.038
```

The selected `lambda.min` for `caret` is 65.4848083, while for `glmnet` is 69.5763174. The `lambda.1se` for `caret` is 942.451867, while for `glmnet` is 1276.037882.

These discrepancies are expected because `caret` and `glmnet` implement cross-validation differently. Although the same `lambda` grid was used, the two functions generate folds independently and compute cross-validated errors and their variability using different internal procedures. Since both `lambda.min` and `lambda.1se` depend on the estimated cross-validation error curve and its standard error, even small differences in fold assignment and error aggregation can lead to different selected tuning parameters.